

CNN_MNIST

August 12, 2019

```
In [6]: from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.layers import Activation
from keras.layers import regularizers
from keras.layers import BatchNormalization
from prettytable import PrettyTable

In [2]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
```

```

print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

In [3]: # convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

In [4]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

WARNING: Logging before flag parsing goes to stderr.

W0811 17:59:04.471589 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages:

W0811 17:59:04.487215 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages:

W0811 17:59:04.487215 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages:

W0811 17:59:04.643455 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages:

W0811 17:59:04.643455 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages:

W0811 17:59:04.659076 7896 deprecation.py:506] From C:\Users\hp\Anaconda3\lib\site-packages\keras:

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

W0811 17:59:04.752821 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.dropout (from tensorflow.nn.dropout_v2) is deprecated and will be removed in a future version.

W0811 17:59:04.768444 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.dropout (from tensorflow.nn.dropout_v2) is deprecated and will be removed in a future version.

W0811 17:59:04.877831 7896 deprecation.py:323] From C:\Users\hp\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.dropout (from tensorflow.nn.dropout_v2) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 113s 2ms/step - loss: 0.2674 - acc: 0.9171 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 2/12

60000/60000 [=====] - 111s 2ms/step - loss: 0.0871 - acc: 0.9742 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 3/12

60000/60000 [=====] - 111s 2ms/step - loss: 0.0658 - acc: 0.9798 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 4/12

60000/60000 [=====] - 109s 2ms/step - loss: 0.0521 - acc: 0.9847 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 5/12

60000/60000 [=====] - 111s 2ms/step - loss: 0.0454 - acc: 0.9863 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 6/12

60000/60000 [=====] - 111s 2ms/step - loss: 0.0418 - acc: 0.9874 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 7/12

60000/60000 [=====] - 112s 2ms/step - loss: 0.0366 - acc: 0.9887 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 8/12

60000/60000 [=====] - 113s 2ms/step - loss: 0.0325 - acc: 0.9899 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 9/12

60000/60000 [=====] - 113s 2ms/step - loss: 0.0302 - acc: 0.9903 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 10/12

60000/60000 [=====] - 112s 2ms/step - loss: 0.0273 - acc: 0.9917 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 11/12

60000/60000 [=====] - 113s 2ms/step - loss: 0.0279 - acc: 0.9915 - val_loss: 0.0283 - val_acc: 0.9924

Epoch 12/12

60000/60000 [=====] - 113s 2ms/step - loss: 0.0250 - acc: 0.9920 - val_loss: 0.0283 - val_acc: 0.9924

Test loss: 0.028258230628092677

Test accuracy: 0.9924

1 Assignment

We have to make about 3 distinct architectures of convolutional networks.

1.0.1 Architecture 1:

```
In [7]: model_1 = Sequential()
        model_1.add(Conv2D(32, kernel_size=(3, 3),
```

```

        activation='relu',
        input_shape=input_shape))
model_1.add(Conv2D(32, (3, 3), activation='relu'))
model_1.add(MaxPooling2D(pool_size=(3, 3)))
model_1.add(BatchNormalization())
model_1.add(Dropout(0.5))
model_1.add(Flatten())
model_1.add(Dense(64, activation='relu'))
model_1.add(Dense(num_classes, activation='softmax'))

model_1.compile(loss=keras.losses.categorical_crossentropy,
                optimizer=keras.optimizers.Adadelta(),
                metrics=['accuracy'])

model_1.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_1 = model_1.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_1[0])
print('Test accuracy:', score_1[1])

```

W0811 18:35:14.560059 7896 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 73s 1ms/step - loss: 0.1829 - acc: 0.9419 - val.
Epoch 2/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0683 - acc: 0.9788 - val.
Epoch 3/12
60000/60000 [=====] - 89s 1ms/step - loss: 0.0523 - acc: 0.9841 - val.
Epoch 4/12
60000/60000 [=====] - 78s 1ms/step - loss: 0.0447 - acc: 0.9857 - val.
Epoch 5/12
60000/60000 [=====] - 79s 1ms/step - loss: 0.0367 - acc: 0.9883 - val.
Epoch 6/12
60000/60000 [=====] - 77s 1ms/step - loss: 0.0325 - acc: 0.9894 - val.
Epoch 7/12
60000/60000 [=====] - 77s 1ms/step - loss: 0.0302 - acc: 0.9902 - val.
Epoch 8/12
60000/60000 [=====] - 76s 1ms/step - loss: 0.0279 - acc: 0.9908 - val.
Epoch 9/12
60000/60000 [=====] - 76s 1ms/step - loss: 0.0244 - acc: 0.9922 - val.
Epoch 10/12
60000/60000 [=====] - 76s 1ms/step - loss: 0.0221 - acc: 0.9928 - val.

```

```

Epoch 11/12
60000/60000 [=====] - 77s 1ms/step - loss: 0.0198 - acc: 0.9940 - val
Epoch 12/12
60000/60000 [=====] - 76s 1ms/step - loss: 0.0205 - acc: 0.9933 - val
Test loss: 0.028312752436022673
Test accuracy: 0.9919

```

1.0.2 Architecture 2

```

In [9]: model_2 = Sequential()
        model_2.add(Conv2D(32, kernel_size=(5, 5),
                           activation='relu',
                           input_shape=input_shape))
        model_2.add(Conv2D(64, (3, 3), activation='relu'))
        model_2.add(BatchNormalization())
        model_1.add(Dropout(0.1))
        model_2.add(Conv2D(64, (2, 2), activation='relu'))
        model_2.add(MaxPooling2D(pool_size=(3, 3)))
        model_2.add(BatchNormalization())
        model_2.add(Flatten())
        model_2.add(Dense(64, activation='relu'))
        model_2.add(BatchNormalization())
        model_2.add(Dropout(0.25))
        model_2.add(Dense(num_classes, activation='softmax'))

        model_2.compile(loss=keras.losses.categorical_crossentropy,
                        optimizer=keras.optimizers.Adam(),
                        metrics=['accuracy'])

        model_2.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
        score_2 = model_2.evaluate(x_test, y_test, verbose=0)
        print('Test loss:', score_2[0])
        print('Test accuracy:', score_2[1])

Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 303s 5ms/step - loss: 0.1254 - acc: 0.9651 - va
Epoch 2/12
60000/60000 [=====] - 303s 5ms/step - loss: 0.0431 - acc: 0.9879 - va
Epoch 3/12
60000/60000 [=====] - 302s 5ms/step - loss: 0.0320 - acc: 0.9904 - va
Epoch 4/12
60000/60000 [=====] - 300s 5ms/step - loss: 0.0230 - acc: 0.9928 - va

```

```

Epoch 5/12
60000/60000 [=====] - 305s 5ms/step - loss: 0.0198 - acc: 0.9941 - va
Epoch 6/12
60000/60000 [=====] - 328s 5ms/step - loss: 0.0168 - acc: 0.9945 - va
Epoch 7/12
60000/60000 [=====] - 332s 6ms/step - loss: 0.0143 - acc: 0.9953 - va
Epoch 8/12
60000/60000 [=====] - 317s 5ms/step - loss: 0.0134 - acc: 0.9957 - va
Epoch 9/12
60000/60000 [=====] - 325s 5ms/step - loss: 0.0105 - acc: 0.9966 - va
Epoch 10/12
60000/60000 [=====] - 311s 5ms/step - loss: 0.0097 - acc: 0.9969 - va
Epoch 11/12
60000/60000 [=====] - 311s 5ms/step - loss: 0.0101 - acc: 0.9966 - va
Epoch 12/12
60000/60000 [=====] - 315s 5ms/step - loss: 0.0077 - acc: 0.9977 - va
Test loss: 0.03223905730191964
Test accuracy: 0.9913

```

1.0.3 Architecture 3:

```

In [10]: model_3 = Sequential()
          model_3.add(Conv2D(32, kernel_size=(2, 2),
                             activation='sigmoid',
                             input_shape=input_shape))
          model_3.add(Conv2D(64, (3, 3), activation='sigmoid'))
          model_3.add(BatchNormalization())
          model_3.add(MaxPooling2D(pool_size=(3, 3)))
          model_3.add(Dropout(0.5))
          model_3.add(Flatten())
          model_3.add(Dense(128, activation='sigmoid'))
          model_3.add(BatchNormalization())
          model_3.add(Dropout(0.2))
          model_3.add(Dense(num_classes, activation='softmax'))

          model_3.compile(loss=keras.losses.categorical_crossentropy,
                          optimizer=keras.optimizers.Adagrad(),
                          metrics=['accuracy'])

          model_3.fit(x_train, y_train,
                      batch_size=batch_size,
                      epochs=epochs,
                      verbose=1,
                      validation_data=(x_test, y_test))
          score_3 = model_3.evaluate(x_test, y_test, verbose=0)
          print('Test loss:', score_3[0])
          print('Test accuracy:', score_3[1])

```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/12
60000/60000 [=====] - 316s 5ms/step - loss: 0.1750 - acc: 0.9469 - va
Epoch 2/12
60000/60000 [=====] - 319s 5ms/step - loss: 0.0886 - acc: 0.9730 - va
Epoch 3/12
60000/60000 [=====] - 317s 5ms/step - loss: 0.0738 - acc: 0.9772 - va
Epoch 4/12
60000/60000 [=====] - 327s 5ms/step - loss: 0.0647 - acc: 0.9802 - va
Epoch 5/12
60000/60000 [=====] - 326s 5ms/step - loss: 0.0578 - acc: 0.9822 - va
Epoch 6/12
60000/60000 [=====] - 324s 5ms/step - loss: 0.0552 - acc: 0.9826 - va
Epoch 7/12
60000/60000 [=====] - 324s 5ms/step - loss: 0.0515 - acc: 0.9846 - va
Epoch 8/12
60000/60000 [=====] - 323s 5ms/step - loss: 0.0474 - acc: 0.9854 - va
Epoch 9/12
60000/60000 [=====] - 323s 5ms/step - loss: 0.0452 - acc: 0.9857 - va
Epoch 10/12
60000/60000 [=====] - 319s 5ms/step - loss: 0.0427 - acc: 0.9868 - va
Epoch 11/12
60000/60000 [=====] - 318s 5ms/step - loss: 0.0407 - acc: 0.9878 - va
Epoch 12/12
60000/60000 [=====] - 332s 6ms/step - loss: 0.0381 - acc: 0.9885 - va
Test loss: 0.03459447893754113
Test accuracy: 0.9881
```

1.0.4 Architecture 4:

```
In [5]: # Activation = selu & softmax , without Batch normalization and dropout, Initializer =
# bias_initializer, Optimizer = Adagrad
epochs = 10
batch_size = 256

model_4 = Sequential()
model_4.add(Conv2D(32, kernel_size=(5, 5),
                  activation='tanh',
                  input_shape=input_shape))
model_4.add(Dense(45,
                  kernel_initializer='random_uniform'))
model_4.add(Conv2D(64, (4, 4), activation='softmax'))
model_4.add(BatchNormalization())
model_4.add(Dropout(0.25))
model_4.add(MaxPooling2D(pool_size=(3, 3)))
model_4.add(Flatten())
model_4.add(Dense(128, activation='sigmoid'))
```

```

model_4.add(BatchNormalization())
model_4.add(Dropout(0.5))
model_4.add(Dense(num_classes, activation='softmax'))

model_4.compile(loss=keras.losses.categorical_crossentropy,
                 optimizer=keras.optimizers.Adamax(),
                 metrics=['accuracy'])

model_4.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_4 = model_4.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_4[0])
print('Test accuracy:', score_4[1])

```

W0811 22:37:34.260540 6576 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

W0811 22:37:34.260540 6576 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

W0811 22:37:34.338661 6576 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

W0811 22:37:34.416781 6576 deprecation.py:506] From C:\Users\hp\Anaconda3\lib\site-packages\k
Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

W0811 22:37:34.432404 6576 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

W0811 22:37:34.573019 6576 deprecation_wrapper.py:119] From C:\Users\hp\Anaconda3\lib\site-pa

W0811 22:37:34.682407 6576 deprecation.py:323] From C:\Users\hp\Anaconda3\lib\site-packages\t
Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 354s 6ms/step - loss: 0.1780 - acc: 0.9479 - va

Epoch 2/10

60000/60000 [=====] - 378s 6ms/step - loss: 0.0719 - acc: 0.9778 - va

Epoch 3/10

60000/60000 [=====] - 360s 6ms/step - loss: 0.0588 - acc: 0.9825 - va

Epoch 4/10

60000/60000 [=====] - 368s 6ms/step - loss: 0.0494 - acc: 0.9842 - va

Epoch 5/10

60000/60000 [=====] - 366s 6ms/step - loss: 0.0438 - acc: 0.9871 - va

Epoch 6/10

60000/60000 [=====] - 362s 6ms/step - loss: 0.0389 - acc: 0.9884 - va


```

Epoch 7/10
60000/60000 [=====] - 397s 7ms/step - loss: 0.0340 - acc: 0.9893 - va
Epoch 8/10
60000/60000 [=====] - 404s 7ms/step - loss: 0.0320 - acc: 0.9901 - va
Epoch 9/10
60000/60000 [=====] - 387s 6ms/step - loss: 0.0275 - acc: 0.9914 - va
Epoch 10/10
60000/60000 [=====] - 381s 6ms/step - loss: 0.0259 - acc: 0.9918 - va
Test loss: 0.036352840721810935
Test accuracy: 0.9884

```

1.1 Conclusions

```

In [11]: number= [1,2,3,4]
         model = ["Architecture 1","Architecture 2","Architecture 3","Architecture 4"]
         opt   = ['Adadelata', 'Adam', 'Adagrad', 'Adamax']
         act   = ['relu', 'relu', 'sigmoid', 'tanh,sigmoid']
         loss  = [0.028312752436022673,0.03223905730191964,0.03459447893754113,score_4[0]]
         acc   = [0.9919,0.9913,0.9881,score_4[1]]

```

```

#Initialize Prettytable
pt = PrettyTable()
pt.add_column("Sr.No.", number)
pt.add_column("Model", model)
pt.add_column("Optimizer", opt)
pt.add_column('Activation ', act)
pt.add_column("Test Loss", loss)
pt.add_column("Test Accuracy", acc)
print(pt)

```

Sr.No.	Model	Optimizer	Activation	Test Loss	Test Accuracy
1	Architecture 1	Adadelata	relu	0.028312752436022673	0.9919
2	Architecture 2	Adam	relu	0.03223905730191964	0.9913
3	Architecture 3	Adagrad	sigmoid	0.03459447893754113	0.9881
4	Architecture 4	Adamax	tanh,sigmoid	0.036352840721810935	0.9884

Architecture 1 is getting highest accuracy of 99.19% with minimum loss, in which I have used Adadelata optimizer and relu activation.

1.1.1 I have tried 4 CNN architectures:

1.1.2 1. with different number of layers

1.1.3 2. different kernel values: (2,2), (3,3), (5,5)

1.1.4 3. different activations: relu, tanh, sigmoid

1.1.5 4. different optimizers: adam, adadelta, adagrad, adamax

1.1.6 5. adding batchnormalisation layers

1.1.7 6. adding some dropout layers, maxpool layers, flattening layers and dense layers