

04 Amazon Fine Food Reviews Analysis_NaiveBayes

June 14, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [86]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from bs4 import BeautifulSoup
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import MultinomialNB
from prettytable import PrettyTable
import math
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000 """)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)

def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (100000, 10)

```

Out[2]:

```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	1	1303862400	
1	0	0	0	1346976000	
2	1	1	1	1219017600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews

```

```
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()
```

```
(80668, 7)
```

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head(3)
```

```

Out [7]:      Id  ProductId      UserId      ProfileName  HelpfulnessNumerator  \
0   78445  B000HDL1RQ  AR5J8UI46CURR  Geetha Krishnan                2
1  138317  B000HDOPYC  AR5J8UI46CURR  Geetha Krishnan                2
2  138277  B000HDOPYM  AR5J8UI46CURR  Geetha Krishnan                2

      HelpfulnessDenominator  Score      Time  \
0                        2      5  1199577600
1                        2      5  1199577600
2                        2      5  1199577600

                        Summary  \
0  LOACKER QUADRATINI VANILLA WAFERS
1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS

                        Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
        sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

In [9]: #Deduplication of entries
        final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
        final.shape

Out [9]: (87775, 10)

In [10]: #Checking to see how much % of data still remains
         (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100

Out [10]: 87.775

```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```

In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()

Out[11]:
   Id  ProductId  UserId  ProfileName \
0  64422  B000MIDR0Q  A161DK06JJMCYF  J. E. Stephens "Jeanne"
1  44737  B001EQ55RW  A2V0I904FH7ABY                      Ram

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      3                      1      5  1224892800
1                      3                      2      4  1212883200

   Summary \
0          Bought This for My Son at College
1  Pure cocoa taste with crunchy almonds inside

   Text
0  My son loves spaghetti so I didn't hesitate or...
1  It was almost a 'love at first bite' - the per...

In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(87773, 10)

Out[13]:
1    73592
0    14181
Name: Score, dtype: int64

```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags

2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [0]: # printing some random reviews
        sent_0 = final['Text'].values[0]
        print(sent_0)
        print("="*50)

        sent_1000 = final['Text'].values[1000]
        print(sent_1000)
        print("="*50)

        sent_1500 = final['Text'].values[1500]
        print(sent_1500)
        print("="*50)

        sent_4900 = final['Text'].values[4900]
        print(sent_4900)
        print("="*50)
```

```
Why is this $[...] when the same product is available for $[...] here?<br />http://www.amazon.
=====
I recently tried this flavor/brand and was surprised at how delicious these chips are.  The be
=====
Wow.  So far, two two-star reviews.  One obviously had no idea what they were ordering; the otl
=====
love to order my coffee on amazon.  easy and shows up quickly.<br />This k cup is great coffee
=====
```

```
In [0]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
        sent_0 = re.sub(r"http\S+", "", sent_0)
        sent_1000 = re.sub(r"http\S+", "", sent_1000)
        sent_150 = re.sub(r"http\S+", "", sent_1500)
        sent_4900 = re.sub(r"http\S+", "", sent_4900)

        print(sent_0)
```

```
Why is this $[...] when the same product is available for $[...] here?<br /> /><br />The Victor
```

```
In [0]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-
        # from bs4 import BeautifulSoup
```

```

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

Why is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The be

=====

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the otl

=====

love to order my coffee on amazon. easy and shows up quickly.This k cup is great coffee. dca

In [14]: # <https://stackoverflow.com/a/47091490/4084039>

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase

```

In [0]: sent_1500 = decontracted(sent_1500)


```
print(sent_1500)
print("="*50)
```

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other was
=====

```
In [0]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />
The Victor

```
In [0]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Wow So far two two star reviews One obviously had no idea what they were ordering the other was

```
In [15]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                'you'll', 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n't',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [16]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
```

```

sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

100%|| 87773/87773 [00:35<00:00, 2469.01it/s]

In [17]: preprocessed_reviews[1500]

Out[17]: 'way hot blood took bite jig lol'

5 [4] Splitting the Data

Before splitting the data in train, cross validation and test , we put all the pre processed reviews in x and corresponding class labels into y.

```

In [18]: x = preprocessed_reviews
         y = final["Score"].values

```

Splitting the data as train data, cross validation data and test data

```

In [19]: # splitting the data into 3 parts for further process,
         # train data, cross validation data and test data

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30) # t
x_train, x_cv, y_train, y_cv = train_test_split(x_train, y_train, test_size=0.30) # t

In [20]: # number of rows in each data set, train, cross validation and test data respectively
print(len(x_train))
print(len(x_cv))
print(len(x_test))

```

43008

18433

26332

6 [5] Featurization

6.1 [5.1] BAG OF WORDS

```

In [24]: # BoW using scikit-learn

```

```

count_vect = CountVectorizer(min_df=10, max_features=10000)
count_vect.fit(x_train)

print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

# we use the fitted CountVectorizer to convert the text to vector
x_train_bow = count_vect.transform(x_train)
x_cv_bow     = count_vect.transform(x_cv)
x_test_bow   = count_vect.transform(x_test)

print("After vectorizations")

print(x_train_bow.shape, y_train.shape)
print(x_cv_bow.shape, y_cv.shape)
print(x_test_bow.shape, y_test.shape)

some feature names ['ability', 'able', 'absence', 'absent', 'absolute', 'absolutely', 'absolu
=====
After vectorizations
(43008, 8076) (43008,)
(18433, 8076) (18433,)
(26332, 8076) (26332,)

```

6.2 [5.2] TF-IDF

In [23]: *# TFIDF using scikit-learn*

```

tf_idf_vect = TfidfVectorizer(min_df=10, max_features=10000) #in scikit
tf_idf_vect.fit(x_train)

print("some sample features",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

# we use fit() method to learn the vocabulary from x_train
# and now transform text data to vectors using transform() method

x_train_tfidf = tf_idf_vect.transform(x_train)
x_cv_tfidf     = tf_idf_vect.transform(x_cv)
x_test_tfidf   = tf_idf_vect.transform(x_test)

print("After featurization\n")

print(x_train_tfidf.shape, y_train.shape)
print(x_cv_tfidf.shape, y_cv.shape)
print(x_test_tfidf.shape, y_test.shape)

some sample features ['ability', 'able', 'absence', 'absent', 'absolute', 'absolutely', 'absolu

```

=====

After featurization

(43008, 8076) (43008,)
(18433, 8076) (18433,)
(26332, 8076) (26332,)

7 [6] Assignment 4: Apply Naive Bayes

```
<li><strong>Apply Multinomial NaiveBayes on these feature sets</strong>
  <ul>
    <li><font color='red'>SET 1:</font>Review text, preprocessed one converted into vectors</li>
    <li><font color='red'>SET 2:</font>Review text, preprocessed one converted into vectors</li>
  </ul>
</li>
<br>
<li><strong>The hyper paramter tuning(find best Alpha)</strong>
  <ul>
<li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom.com'>Applied AI</a>
<li>Consider a wide range of alpha values for hyperparameter tuning, start as low as 0.00001</li>
<li>Find the best hyper paramter using k-fold cross validation or simple cross validation data</li>
<li>Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task</li>
  </ul>
</li>
<br>
<li><strong>Feature importance</strong>
  <ul>
<li>Find the top 10 features of positive class and top 10 features of negative class for both classes</li>
  </ul>
</li>
<br>
<li><strong>Feature engineering</strong>
  <ul>
<li>To increase the performance of your model, you can also experiment with with feature engineering</li>
    <ul>
      <li>Taking length of reviews as another feature.</li>
      <li>Considering some features from review summary as well.</li>
    </ul>
  </ul>
</li>
<br>
<li><strong>Representation of results</strong>
  <ul>
<li>You need to plot the performance of model both on train data and cross validation data for both classes</li>
<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and find the performance on test data</li>
  </ul>
</li>
```

```

<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='confusion_matrix.png' width=300px></li>
</ul>
</li>
<br>
<li><strong>Conclusion</strong>
<ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
<img src='summary.JPG' width=400px>
</li>
</ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

8 Applying Multinomial Naive Bayes

8.1 [6.1] Applying Naive Bayes on BOW, SET 1

So let's find a best value for our hyperparameter here using GridSearch hyperparameter tuning.

```

In [44]: lg_alp = []
         alpha = [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]

         for i in alpha :
             lg_alp.append(math.log10(i))

         print(lg_alp)

```

```
[-4.0, -3.0, -2.0, -1.0, 0.0, 1.0, 2.0, 3.0, 4.0]
```

```

In [40]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV

         parameters = {'alpha':[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]}
         nb = MultinomialNB()

         grid = GridSearchCV(nb, parameters, cv=4, scoring='roc_auc', n_jobs=-1)
         grid.fit(x_train_bow, y_train)

         train_auc_bow      = grid.cv_results_['mean_train_score']
         cv_auc_bow         = grid.cv_results_['mean_test_score']

```


```

plt.plot(lg_alp, train_auc_bow, label='Train AUC')
plt.scatter(lg_alp, train_auc_bow, label='Train AUC')

plt.plot(lg_alp, cv_auc_bow, label='CV AUC')
plt.scatter(lg_alp, cv_auc_bow, label='CV AUC')

plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
print("Model with best parameters :\n",grid.best_estimator_)

```



```

Model with best parameters :
MultinomialNB(alpha=1, class_prior=None, fit_prior=True)

```

```

In [41]: best_alpha = 1

```

```


In [77]: nb = MultinomialNB(alpha = best_alpha)
nb.fit(x_train_bow, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of

```

```
# not the predicted outputs
```


```
train_fpr_bow, train_tpr_bow, thresholds_bow = roc_curve(y_train, nb.predict_proba(x_train_bow))  
test_fpr_bow, test_tpr_bow, thresholds_bow = roc_curve(y_test, nb.predict_proba(x_test_bow))
```

```
plt.plot(train_fpr_bow, train_tpr_bow, label="train AUC =" + str(auc(train_fpr_bow, train_tpr_bow)))  
plt.plot(test_fpr_bow, test_tpr_bow, label="test AUC =" + str(auc(test_fpr_bow, test_tpr_bow)))  
plt.legend()  
plt.xlabel("alpha: hyperparameter")  
plt.ylabel("AUC")  
plt.title("ERROR PLOTS")  
plt.show()
```



```
In [120]: print("Train Confusion Matrix")  
cm_tr = confusion_matrix(y_train, grid.predict(x_train_bow))  
sns.heatmap(cm_tr, annot=True, fmt='d') # annot=True to annotate cells  
  
# labels, title and ticks  
plt.xlabel('Predicted labels')  
plt.ylabel('True labels')  
plt.title('Train Confusion Matrix')  
plt.show()
```


Train Confusion Matrix



```
In [121]: print("Test Confusion Matrix")
          cm_te = confusion_matrix(y_test, grid.predict(x_test_bow))
          sns.heatmap(cm_te, annot=True, fmt='g')

          # labels, title and ticks
          plt.xlabel('Predicted labels')
          plt.ylabel('True labels')
          plt.title('Train Confusion Matrix')
          plt.show()
```

Test Confusion Matrix



8.1.1 [6.1.1] Top 10 important features of positive class from SET 1

In [82]: `nb.classes_`

Out[82]: `array([0, 1], dtype=int64)`

In [83]: *# Now we can find log probabilities of different features for both the classes*
`class_features = nb.feature_log_prob_`

row_0 is for 'negative' class and row_1 is for 'positive' class

`negative_features = class_features[0]`

`positive_features = class_features[1]`

Getting all feature names

`feature_names = count_vect.get_feature_names()`

Sorting 'negative_features' and 'positive_features' in descending order using argsort

`sorted_negative_features = np.argsort(negative_features)[::-1]`

`sorted_positive_features = np.argsort(positive_features)[::-1]`

`print("\n\nTop 10 Important Features and their log probabilities For Positive Class :")`

`for i in list(sorted_positive_features[0:10]):`

`print("%s\t -->\t%f" %(feature_names[i],positive_features[i]))`

Top 10 Important Features and their log probabilities For Positive Class :

not	-->	-3.679396
like	-->	-4.480949
good	-->	-4.621951
great	-->	-4.676597
one	-->	-4.832851
taste	-->	-4.916672
coffee	-->	-4.930744
flavor	-->	-5.005918
would	-->	-5.020401
love	-->	-5.028871

8.1.2 [6.1.2] Top 10 important features of negative class from SET 1

```
In [84]: print("Top 10 Important Features and their log probabilities For Negative Class :\n\n")
        for i in list(sorted_negative_features[0:10]):
            print("%s\t -->\t%f" %(feature_names[i],negative_features[i]))
```

Top 10 Important Features and their log probabilities For Negative Class :

not	-->	-3.235728
like	-->	-4.350658
would	-->	-4.637823
taste	-->	-4.649252
product	-->	-4.691221
one	-->	-4.871003
good	-->	-5.062585
coffee	-->	-5.081188
flavor	-->	-5.137463
no	-->	-5.181554

```
In [114]: # evaluating train: accuracy, precison, recall and f1_score
```

```
predictions = nb.predict(x_train_bow)

acc_tr = accuracy_score(y_train, predictions)
print('Accuracy score: ', acc_tr)

pre_tr = precision_score(y_train, predictions)
print('Precision score: ', pre_tr)
```

```

rec_tr = recall_score(y_train, predictions)
print('Recall score: ', rec_tr)

f1_tr = f1_score(y_train, predictions)
print('F1 score: ', f1_tr)

```

```

Accuracy score:  0.9190150669642857
Precision score:  0.9578238575533893
Recall score:    0.9448072700997693
F1 score:        0.9512710382360758

```

```

In [113]: # evaluating test: accuracy, precision, recall and f1_score
predictions = nb.predict(x_test_bow)

```

```

acc_te = accuracy_score(y_test, predictions)
print('Accuracy score: ', acc_te)

pre_te = precision_score(y_test, predictions)
print('Precision score: ', pre_te)

rec_te = recall_score(y_test, predictions)
print('Recall score: ', rec_te)

f1_te = f1_score(y_test, predictions)
print('F1 score: ', f1_te)

```

```

Accuracy score:  0.9011089169071852
Precision score:  0.9430716723549488
Recall score:    0.9387144992526159
F1 score:        0.9408880414056116

```

8.2 [6.2] Applying Naive Bayes on TFIDF, SET 2

Hyperparameter tuning using grid search

```

In [93]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearch

```

```

nb = MultinomialNB()
parameters = {'alpha':[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]}
grid = GridSearchCV(nb, parameters, cv=3, scoring='roc_auc', n_jobs=-1)
grid.fit(x_train_tfidf, y_train)

train_auc_tfidf      = grid.cv_results_['mean_train_score']
cv_auc_tfidf         = grid.cv_results_['mean_test_score']

plt.plot(lg_alp, train_auc_tfidf, label='Train AUC')

```


```

plt.scatter(lg_alp, train_auc_tfidf, label='Train AUC')

plt.plot(lg_alp, cv_auc_tfidf, label='CV AUC')
plt.scatter(lg_alp, cv_auc_tfidf, label='CV AUC')

plt.legend()
plt.xlabel("log(alpha): hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
print("Model with best parameters :\n",grid.best_estimator_)

```



```

Model with best parameters :
MultinomialNB(alpha=0.1, class_prior=None, fit_prior=True)

```

```

In [94]: best_alpha_tfidf = 0.1

```

```

In [96]: nb = MultinomialNB(alpha = best_alpha_tfidf)
nb.fit(x_train_bow, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs

train_fpr_tfidf, train_tpr_tfidf, thresholds_tfidf = roc_curve(y_train, nb.predict_proba(x_train_bow)[:,1])


```

```

test_fpr_tfidf, test_tpr_tfidf, thresholds_tfidf = roc_curve(y_test, nb.predict_proba

plt.plot(train_fpr_tfidf, train_tpr_tfidf, label="train AUC =" + str(auc(train_fpr_tfidf, t
plt.plot(test_fpr_tfidf, test_tpr_tfidf, label="test AUC =" + str(auc(test_fpr_tfidf, t
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```




```

In [117]: print("Train Confusion Matrix")
cm_tr = confusion_matrix(y_train, grid.predict(x_train_tfidf))
sns.heatmap(cm_tr, annot=True, fmt='d') # annot=True to annotate cells

# labels, title and ticks
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Train Confusion Matrix')
plt.show()

```


Train Confusion Matrix



```
In [122]: print("Train Confusion Matrix")
          cm_te = confusion_matrix(y_test, grid.predict(x_test_tfidf))
          sns.heatmap(cm_te, annot=True, fmt= 'd') # annot=True to annotate cells

          # labels, title and ticks
          plt.xlabel('Predicted labels')
          plt.ylabel('True labels')
          plt.title('Train Confusion Matrix')
          plt.show()
```

Train Confusion Matrix



8.2.1 [6.2.1] Top 10 important features of positive class from SET 2

```
In [ ]: # https://github.com/PushpendraSinghChauhan/Amazon-Fine-Food-Reviews/blob/master/Apply
```

```
In [97]: nb.classes_
```

```
Out[97]: array([0, 1], dtype=int64)
```

```
In [98]: # Now we can find log probabilities of different features for both the classes
class_features = nb.feature_log_prob_
```

```
# row_0 is for 'negative' class and row_1 is for 'positive' class
negative_features = class_features[0]
positive_features = class_features[1]
```

```
# Getting all feature names
feature_names = tf_idf_vect.get_feature_names()
```

```
# Sorting 'negative_features' and 'positive_features' in descending order using argsort
sorted_negative_features = np.argsort(negative_features)[::-1]
sorted_positive_features = np.argsort(positive_features)[::-1]
```

```
In [99]: print("\n\nTop 10 Important Features and their log probabilities For Positive Class :")
for i in list(sorted_positive_features[0:10]):
    print("%s\t -->\t%f" % (feature_names[i], positive_features[i]))
```

Top 10 Important Features and their log probabilities For Positive Class :

not	-->	-3.673939
like	-->	-4.475525
good	-->	-4.616536
great	-->	-4.671185
one	-->	-4.827452
taste	-->	-4.911280
coffee	-->	-4.925353
flavor	-->	-5.000535
would	-->	-5.015020
love	-->	-5.023491

8.2.2 [6.2.2] Top 10 important features of negative class from SET 2

```
In [100]: print("Top 10 Important Features and their log probabilities For Negative Class :\n\n")
          for i in list(sorted_negative_features[0:10]):
              print("%s\t -->\t%f" %(feature_names[i],negative_features[i]))
```

Top 10 Important Features and their log probabilities For Negative Class :

not	-->	-3.211281
like	-->	-4.326367
would	-->	-4.613610
taste	-->	-4.625042
product	-->	-4.667025
one	-->	-4.846872
good	-->	-5.038536
coffee	-->	-5.057148
flavor	-->	-5.113451
no	-->	-5.157564

```
In [111]: # evaluating train: accuracy, precison, recall and f1_score
```

```
predictions1 = nb.predict(x_train_tfidf)

acc_tr1 = accuracy_score(y_train, predictions1)
print('Accuracy score: ', acc_tr1)

pre_tr1 = precision_score(y_train, predictions1)
print('Precision score: ', pre_tr1)
```



```

rec_tr1 = recall_score(y_train, predictions1)
print('Recall score: ', rec_tr1)

f1_tr1 = f1_score(y_train, predictions1)
print('F1 score: ', f1_tr1)

```

```

Accuracy score:  0.8858119419642857
Precision score:  0.8829619404456714
Recall score:    0.9954700830947948
F1 score:        0.9358466904416665

```

In [112]: *# evaluating test accuracy, precison, recall, f1_score*

```

predictions1 = nb.predict(x_test_tfidf)
acc_te1 = accuracy_score(y_test, predictions1)
print('Accuracy score: ', acc_te1)

pre_te1 = precision_score(y_test, predictions1)
print('Precision score: ', pre_te1)

rec_te1 = recall_score(y_test, predictions1)
print('Recall score: ', rec_te1)

f1_te1 = f1_score(y_test, predictions1)
print('F1 score: ', f1_te1)

```

```

Accuracy score:  0.873310041014735
Precision score:  0.8712412344994256
Recall score:    0.996104543189745
F1 score:        0.9294982881778604

```

9 [7] Conclusions

In [115]: *# Creating table using PrettyTable library*

```

names = ["MultinomialNB for BoW", "MultinomialNB for TFIDF"]

optimal_alpha = [ best_alpha, best_alpha_tfidf]

#train_acc = [ train_auc_bow, train_auc_tfidf]

train_acc = [ acc_tr, acc_tr1]
test_acc  = [ acc_te, acc_te1]

numbering = [1,2]

```

```

# Initializing prettytable
ptable = PrettyTable()

# Adding columns
ptable.add_column("S.NO.", numbering)
ptable.add_column("MODEL", names)
ptable.add_column("Best Alpha", optimal_alpha)
ptable.add_column("Train Accuracy", train_acc)
ptable.add_column("Test Accuracy", test_acc)
#ptable.add_column("Training Accuracy", train_acc)
#ptable.add_column("Test Accuracy", test_acc)

# Printing the Table
print(ptable)

```

S.NO.	MODEL	Best Alpha	Train Accuracy	Test Accuracy
1	MultinomialNB for BoW	1	0.9190150669642857	0.9011089169071852
2	MultinomialNB for TFIDF	0.1	0.8858119419642857	0.873310041014735